

Towards Trusted Temporal Databases Using Blockchain

By Amin Beirami
MSc student, UOIT

Supervisors: Dr. Ken Pu and Dr. Ying Zhu

October 2018

Trusting data

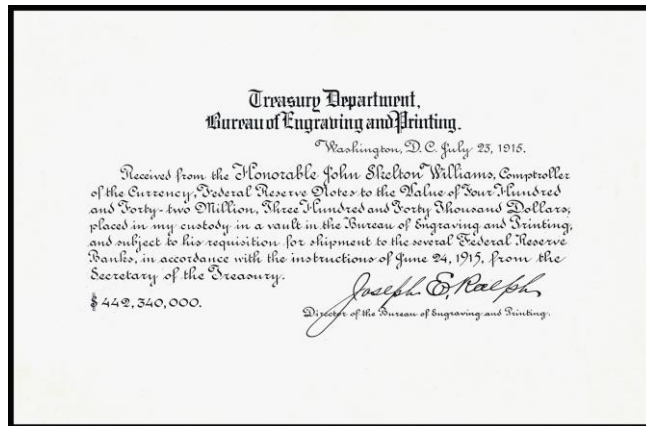
(Big) Data supports important decision making.

It is increasingly important to manage *trust* in
database and data *analytics*

Data is stored as facts in relational databases.

The question is are they trustworthy?

Name	Date	Amount
John Williams	1915-07-23	442,340
Burger Ville	2009-09-15	11.28
Burger Ville	2009-09-15	6.33



(John Williams, 1915-07-23, \$442,340)



(BurgerVille, 2009-09-15, \$11.28)
(BurgerVille, 2009-09-15, \$6.33)

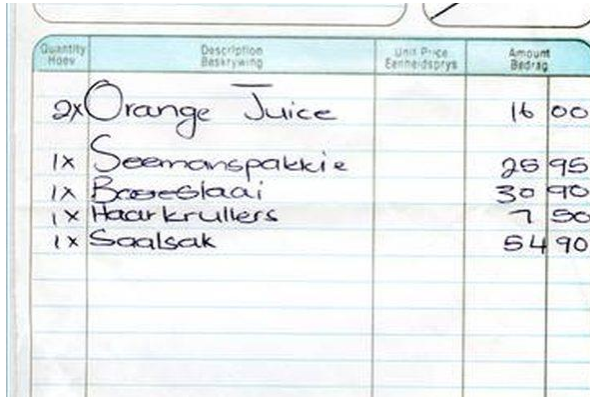
Gaining Trust: an example

Existing methods in establishing (un)trusted data in practice.

Reimbursement for travel expenses.

Gaining Trust: an example

Hand written receipt with *pencil* on *paper*.



A handwritten receipt on lined paper. The receipt is structured as a table with four columns: Quantity (Hoeve), Description (Besrywing), Unit Price (Eenhedsprijs), and Amount (Bedrag). The entries are written in pencil.

Quantity Hoeve	Description Besrywing	Unit Price Eenhedsprijs	Amount Bedrag
2x	Orange Juice		16 00
1x	Seemanspakkie		25 95
1x	Breëslaki		30 90
1x	Haarkrullers		7 50
1x	Soalsak		54 90

Scheme A

Database technology

1. Easy to create
2. Easy to alter (Mutable)
3. Easy to duplicate with forgery
4. Not verifiable

SQLite database file

Gaining Trust: an example

Printed receipt on business ledger

BEST DEAL STORES	
COOKIES	5.00
MILK 65 Fl oz	3.78
SUBTOTAL	8.78
TAX 5%	0.44
TOTAL	9.22
CASH TEND	10.00
CASH DUE	0.78
# ITEMS SOLD 2	
11/03/13 19:53:17	

Scheme B

Database technology

1. Less easy to create
2. Hard to alter (almost immutable)
3. Easy to duplicate with forgery
4. Not verifiable

SQL server with login credentials

Gaining Trust: an example

Printed receipt on business ledger with

1. Signature and Serial number.



Scheme C

Database technology

1. Harder to create
2. Immutable
3. Hard to duplicate with forgery
4. Verifiable (but difficult) using replicated storage

Distributed SQL server (in the cloud)

Gaining Trust: an example

Only business executives can submit
expense claims for reimbursement.

NO RECEIPT

Scheme D

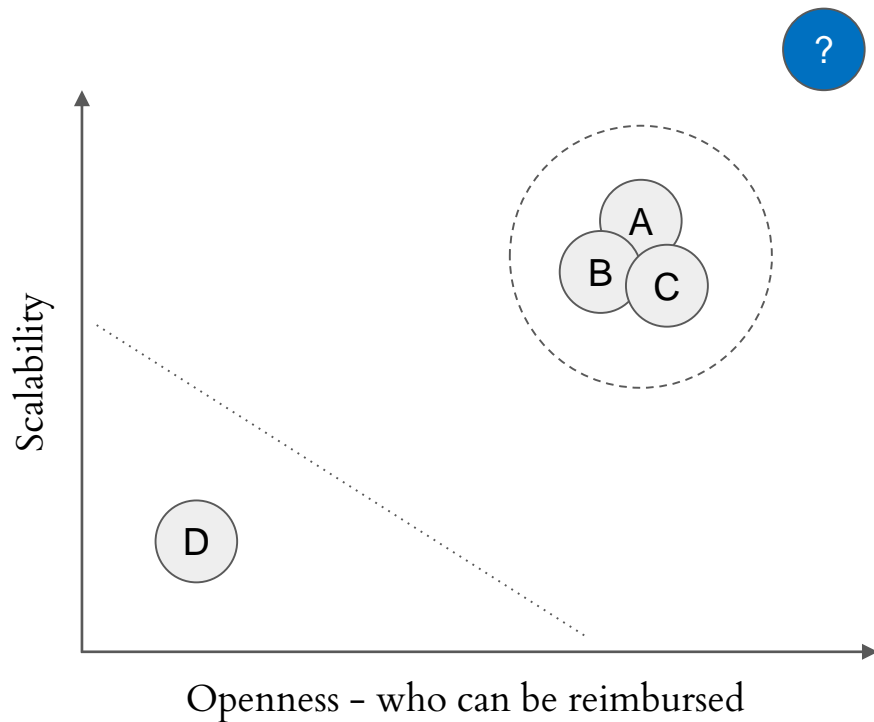
Database technology –

(implicit trust in data)

1. Creation doesn't matter
2. Mutability doesn't matter
3. Difficulty in forgery – N/A
4. Verifiability – N/A

SQL Server with authentication
and access control

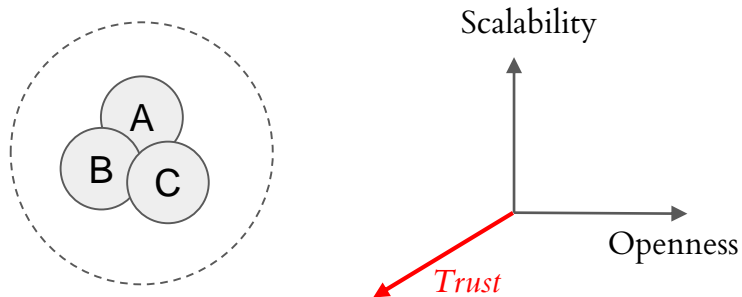
Gaining Trust: an example



Our research interest is

Scalable and Open Trusted Databases

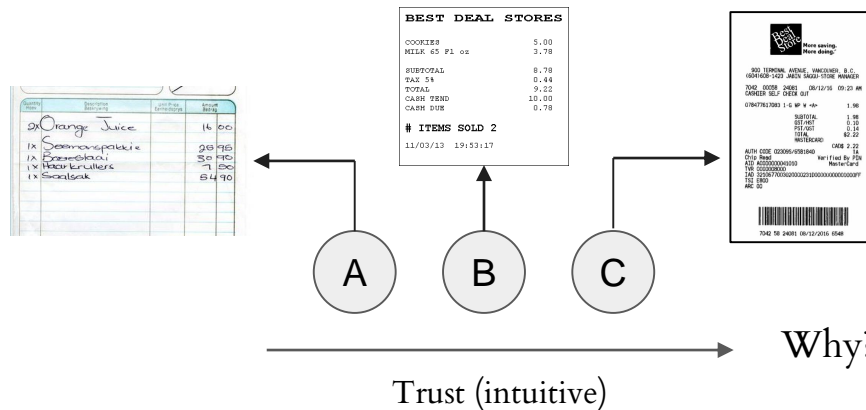
Gaining Trust: an example



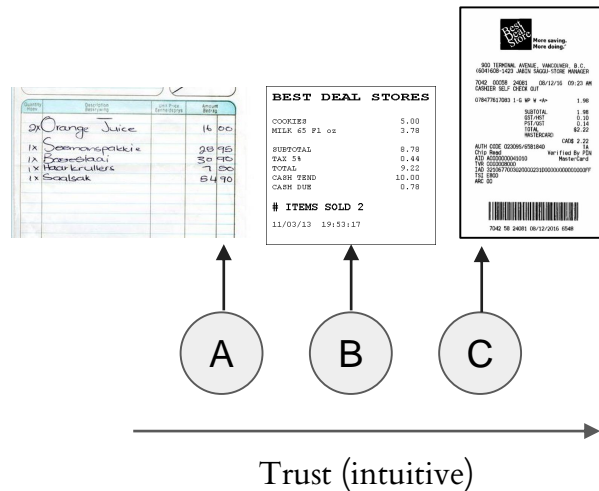
Our research interest is

Scalable and Open

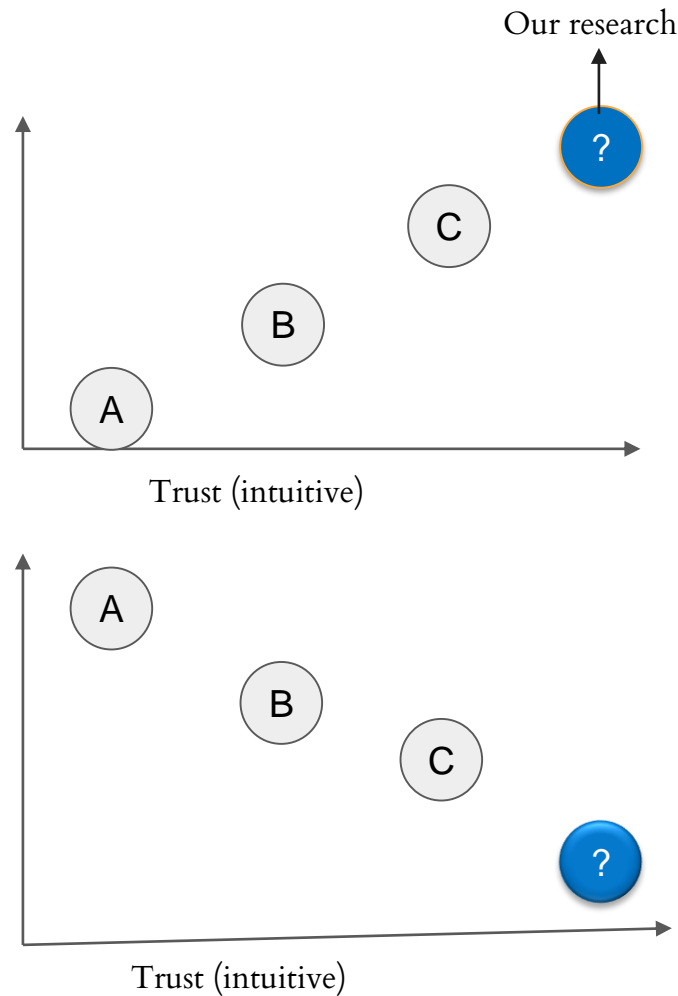
& *Trusted* Databases



Gaining Trust: an example



Computational
complexity in
creating the *perfect*
forgery



Research

A *trusted* transaction logfile to support:

1. Provide auditable data provenance.
2. Temporal relational snapshots at user defined timestamps.

Part 1

Providing trusted data provenance

Logfile to support data provenance

id	Part Name	Manufacturer	Price
010	Console	Nintendo	399
012	Keyboard	Microsoft	89
.	.	.	.
.	.	.	.
.	.	.	.

Data provenance

Why is the *price* of
Nintendo set to \$399?

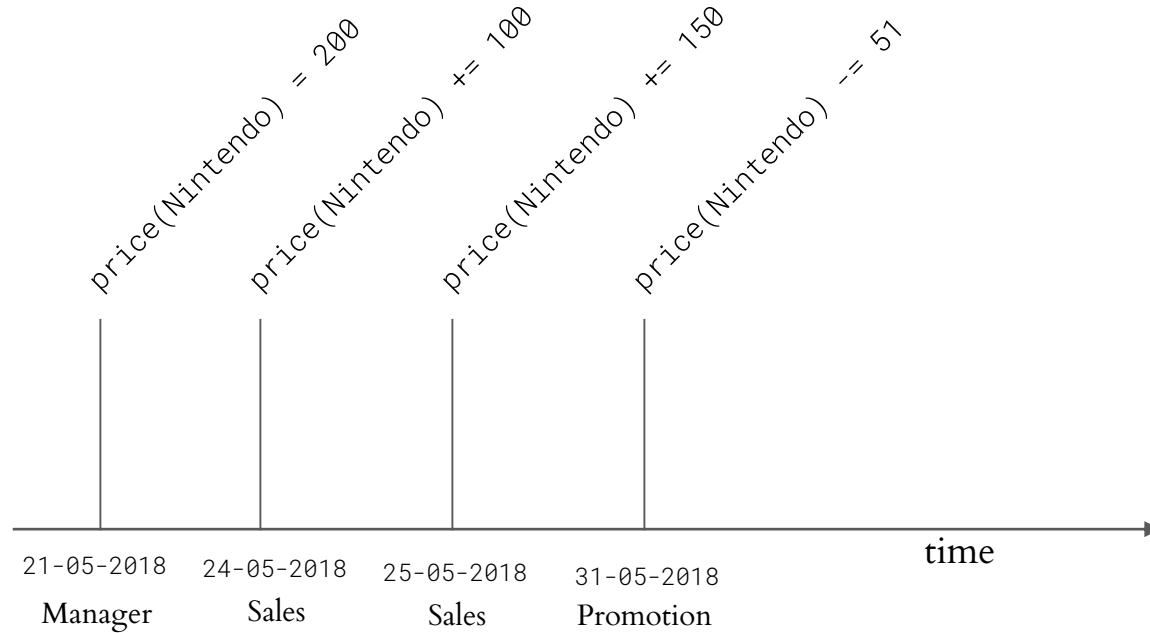
When is the *price* of
Nintendo set to \$399?

Who set the *price* of
Nintendo to \$399?

We need all the historic transactions on the table.

Logfile = **temporal relations**

Logfile to support data provenance



Definition: *temporal relations*

Given a relation r , with attributes $attr(r)$, its temporal relation, written, r^T , is relation with schema:

$$attr(r^T) = attr(r) \cup \{\text{timestamp, deleted}\}$$

Temporal tables

id	Part Name	Manufacturer	Price
010	Console	Nintendo	399
012	Keyboard	Microsoft	89
.	.	.	.
.	.	.	.
.	.	.	.

Relational Table

id	Part Name	Manufacturer	Price	timestamp	deleted
010	Console	Nintendo	200	21-05-2018	False
011	Monitor	Asus	499	23-05-2018	False
012	Keyboard	Microsoft	119	23-05-2018	False
010	Console	Nintendo	300	24-05-2018	False
011	-	-	-	25-05-2018	True
012	Keyboard	Microsoft	89	25-05-2018	False
010	Console	Nintendo	450	27-05-2018	False
010	Console	Nintendo	399	31-05-2018	False
.
.

Temporal Table

Untrusted provenance

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted
010	Console	Nintendo	200	21-05-2018	False
011	Monitor	Asus	499	23-05-2018	False
012	Keyboard	Microsoft	119	23-05-2018	False
010	Console	Nintendo	300	24-05-2018	False
011	-	-	-	25-05-2018	True
012	Keyboard	Microsoft	89	25-05-2018	False
010	Console	Nintendo	450	27-05-2018	False
010	Console	Nintendo	399	31-05-2018	False
.

Who committed the transactions?

How can we trust the transactions?

price(Nintendo) = \$399

Untrusted provenance

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted
010	Console	Nintendo	200	21-05-2018	False
011	Monitor	Asus	499	23-05-2018	False
012	Keyboard	Microsoft	119	23-05-2018	False
010	Console	Nintendo	300	24-05-2018	False
011	-	-	-	25-05-2018	True
012	Keyboard	Microsoft	89	25-05-2018	False
010	Console	Nintendo	450	27-05-2018	False
010	Console	Nintendo	470	31-05-2018	False
.	.	.	399	.	.
.

The temporal relation supports data manipulation expressions.

Adversarial attack by super-user.

We need the data to be immutable and verifiable.

price(Nintendo) = \$399

Temporal database with security information

Given a temporal relation r^T , the trusted version of the relation written, r^{T*} , is relation with schema:

$$\text{attr}(r^{T*}) = \text{attr}(r^T) \cup \{\text{username, signature}\}$$

Digital Signature

Mathematical scheme for presenting the authenticity
of a digital information

Adds verification to the records

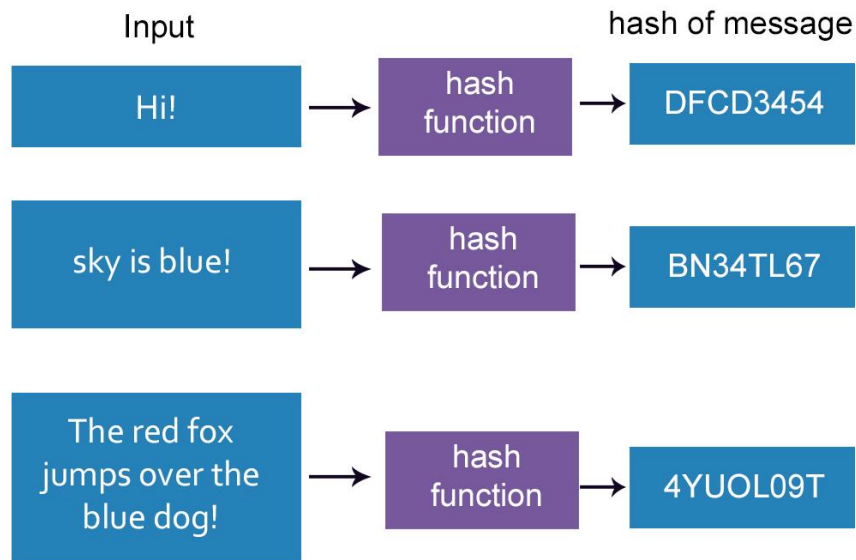
It is unique for user and data

Hash functions

Assume m to be a message of an arbitrary size

$hash(m) \rightarrow sketch$

The function maps m from a larger domain to a **fixed-size string** in a smaller domain



Asymmetric encryption

To encode information such that only authorized parties could read.

Given $\langle K_{private}, K_{public} \rangle \in R^+$
 $K_{private}$ is known to the owner,
 K_{public} is shared with public

$cipher = encrypt(m, K_{public})$

$m = decrypt(cipher, K_{private})$

Function generateKeys(*keySize*):

```
| randomVal ← Random.new()  
| publicKey, privateKey = RSA.generate(keySize, randomVal)  
| return publicKey, privateKey
```

Function encrypt(*m*, *publicKey*):

```
| publicKeyObj = RSA.importKey(publicKey)  
| randomParam ← random.choice()  
| return publicKeyObj.encrypt(m, randomParam)
```

Function decrypt(*enc_message*, *privateKey*):

```
| privateKeyObj = RSA.importKey(privateKey)  
| return privateKeyObj.decrypt(enc_message)
```

Digital signature

$sig = \text{encrypt}(\text{hash}(m), K_{\text{private}})$

$isValid = \text{verify}(m, sig, K_{\text{public}})$

```
Function digitalSignature(m,privateKey):  
    hashVal = hash(m)  
    randomVal = random.choice()  
    privateKeyObj = RSA.importKey(privateKey)  
    return privateKeyObj.encrypt(hashVal,randomVal)
```

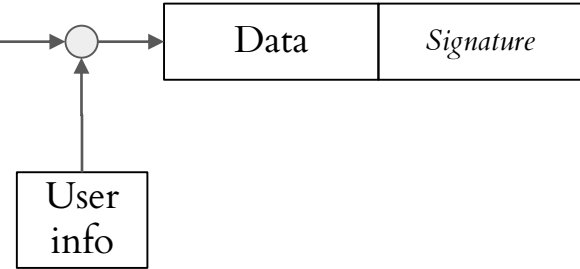
```
Function verifySign(m,signature,publicKey):  
    publicKeyObj = RSA.importKey(publicKey)  
    hashVal = hash(m)  
    signHash = publicKeyObj.decrypt(signature)  
    if hashVal == signHash then  
        | return valid  
    return invalid
```


Signing transactions

Each user is given a unique ID,
and RSA signing key-pairs.

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted
010	Console	Nintendo	200	21-05-2018	False
011	Monitor	Asus	499	23-05-2018	False
012	Keyboard	Microsoft	119	23-05-2018	False
010	Console	Nintendo	300	24-05-2018	False
011	-	-	-	25-05-2018	True
012	Keyboard	Microsoft	89	25-05-2018	False
010	Console	Nintendo	450	27-05-2018	False
010	Console	Nintendo	399	31-05-2018	False
.

*Cryptographic
signing*



Verifiable without revealing
users' secrets

Temporal relations with digital signature

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35
011	-	-	-	25-05-2018	True	manager	09RTF15
010	Console	Nintendo	450	27-05-2018	False	manager	J3IOGLS
012	Keyboard	Microsoft	89	31-05-2018	False	clerk2	78PP00L
.		
.		

We know who committed the transactions

Authenticity of the records are verifiable

price(Nintendo) = \$450

Temporal relations with digital signature

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35
011	-	-	-	25-05-2018	True	manager	09RTF15
010	Console	Nintendo	450	27-05-2018	False	manager	J3IOGLS
010	Console	Nintendo	399	29-05-2018	False	clerk2	SDN0PLT
012	Keyboard	Microsoft	89	31-05-2018	False	clerk2	78PP00L
.

Adversarial attack is more difficult now, but still possible with committing malicious transactions.

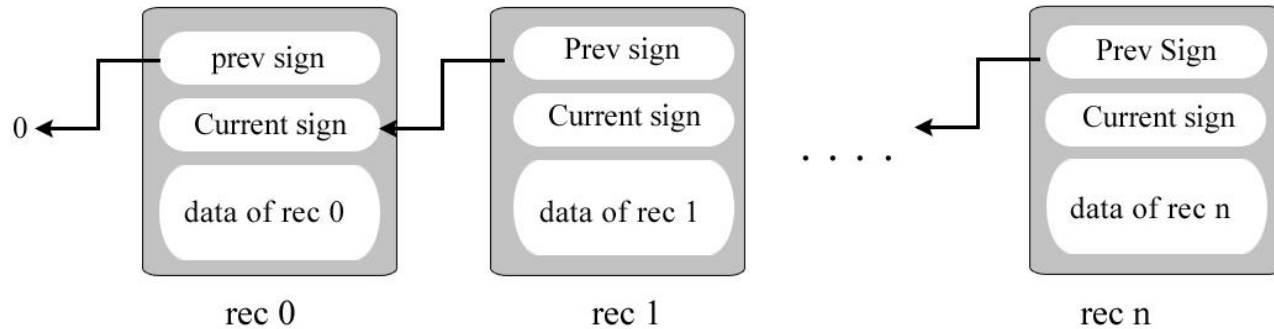
→ *Clerk2 is a superuser*

price(Nintendo) = \$399

Blockchain

Set of blocks chained together by *digital signatures*.

Each block contains its previous block's digital signature



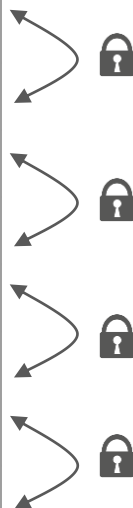
Temporal database with security information

Given a temporal relation with security information r^{T*} , the relation with chained records, written as, r^{T+} has the schema:

$$attr(r^{T+}) = attr(r^{T*}) \cup \{\text{previous_signature}\}$$

Trusted temporal tables with *blockchain*

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35	04P5KM9
011	-	-	-	25-05-2018	True	manager	09RTF15	YUI9T35
012	Keyboard	Microsoft	89	25-05-2018	False	manager	J3IOGLS	09RTF15
010	Console	Nintendo	450	27-05-2018	False	clerk2	78PP00L	J3IOGLS
010	Console	Nintendo	399	31-05-2018	False	manager	ZNTP67W	78PP00L
:	:	:	:	:	:			



Trusted temporal tables with *blockchain*

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35	04P5KM9
011	-	-	100	25-05-2018	True	manager	09RTF15	YUI9T35
012	Keyboard	Microsoft	89	25-05-2018	False	manager	J3IOGLS	09RTF15
010	Console	Nintendo	450	27-05-2018	False	clerk2	78PP00L	J3IOGLS
010	Console	Nintendo	399	31-05-2018	False	manager	ZNTP67W	78PP00L
⋮	⋮	⋮	⋮	⋮	⋮			

*Adversarial attack
by clerk2.*

Trusted temporal tables with *blockchain*

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	100	24-05-2018	False	Clerk 2	OPNJNP3	04P5KM9
011	-	-	-	25-05-2018	True	manager	09RTF15	YUI9T35
012	Keyboard	Microsoft	89	25-05-2018	False	manager	J3IOGLS	09RTF15
010	Console	Nintendo	450	27-05-2018	False	clerk2	78PP00L	J3IOGLS
010	Console	Nintendo	399	31-05-2018	False	manager	ZNTP67W	78PP00L
.			

*Adversarial attack
by clerk2.*

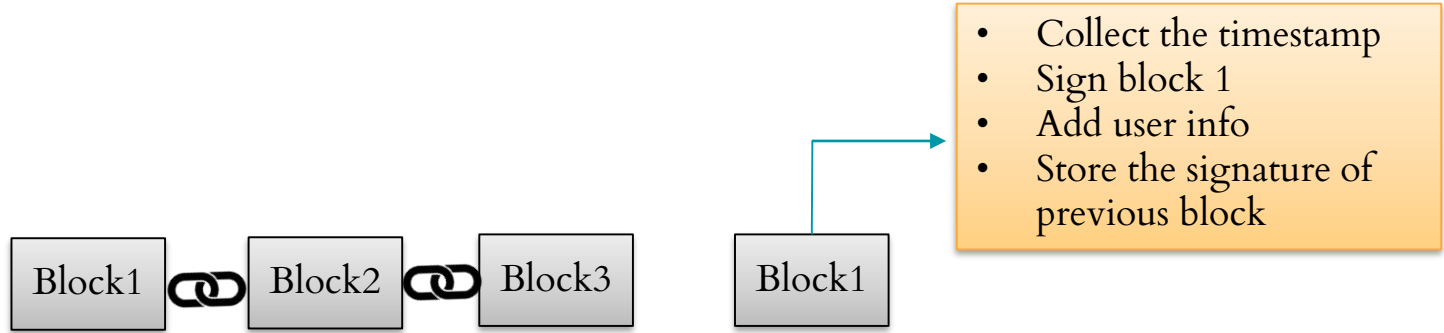
 *Broken chain*

Trusted temporal tables with *blockchain*

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	100	24-05-2018	False	Clerk 2	OPNJNP3	04P5KM9
011	-	-	-	25-05-2018	True	Clerk 2	NYT0903	OPNJNP3
012	Keyboard	Microsoft	89	25-05-2018	False	Clerk 2	YUI89PY	NYT0903
010	Console	Nintendo	450	27-05-2018	False	Clerk 2	09SSTER	YUI89PY
010	Console	Nintendo	399	31-05-2018	False	Clerk 2	IOP7TU8	09SSTER
.			

*Adversarial attack
by clerk2.*

Adding records to the Blockchain



Adding blocks could be done with $O(1)$

Verifying Blockchain

```
Function blockchainVerification(last_id):  
  for  $i \leftarrow 1$  to last_id do  
    publicKey  $\leftarrow$  SELECT publicKey FROM USERS WHERE user = username;  
    validity = verifySign(reci,signature,publicKey)  
    if validity  $\neq$  valid then  
      return Broken chain  
      Break  
    if reci[prev_sgn]  $\neq$  reci-1[signature] then  
      return Broken chain  
      Break  
  end  
  return Valid Chain
```

Part 2

Query Evaluation Over Trusted Temporal Relation

Temporal databases: snapshot creation

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35	04P5KM9
011	-	-	-	25-05-2018	True	manager	09RTF15	YUI9T35
012	Keyboard	Microsoft	89	25-05-2018	False	manager	J3IOGLS	09RTF15
010	Console	Nintendo	450	27-05-2018	False	clerk2	78PP00L	J3IOGLS
010	Console	Nintendo	399	31-05-2018	False	manager	ZNTP67W	78PP00L
.			

What did the relation/record look like in 25-05-2018?

Are the records/ snapshots until then trustworthy?

Temporal databases: snapshot creation

We can generate snapshots using relational database windowing functions

```
snapshot( $r, t$ ) =  
  WITH  $T$  AS (  
    SELECT id, {last_value( $x$ ) as  $x : x \in \text{attr}(r)$ } OVER  $W$   
    FROM  $r^T$   
    WHERE updates  $\leq t$   
    WINDOW  $W$  AS PARTITION BY id ORDER BY updates  
  )  
  SELECT id, { $x : x \in \text{attr}(r)$ } FROM  $T$   
  WHERE NOT  $T.\text{deleted}$ 
```

Temporal databases: snapshot creation

Rec_id	Part Name	Manufacturer	Price	timestamp	deleted	user	sign	pre_sgn
010	Console	Nintendo	200	21-05-2018	False	clerk1	6AE2BTN	0
011	Monitor	Asus	499	23-05-2018	False	manager	NT89RE7	6AE2BTN
012	Keyboard	Microsoft	119	23-05-2018	False	manager	04P5KM9	NT89RE7
010	Console	Nintendo	300	24-05-2018	False	clerk1	YUI9T35	04P5KM9
011	-	-	-	25-05-2018	True	manager	09RTF15	YUI9T35
012	Keyboard	Microsoft	89	25-05-2018	False	manager	J3IOGLS	09RTF15
010	Console	Nintendo	450	27-05-2018	False	clerk2	78PP00L	J3IOGLS
010	Console	Nintendo	399	31-05-2018	False	manager	ZNTP67W	78PP00L
⋮	⋮	⋮	⋮	⋮	⋮			

Records to be processed.

Temporal databases: snapshot creation

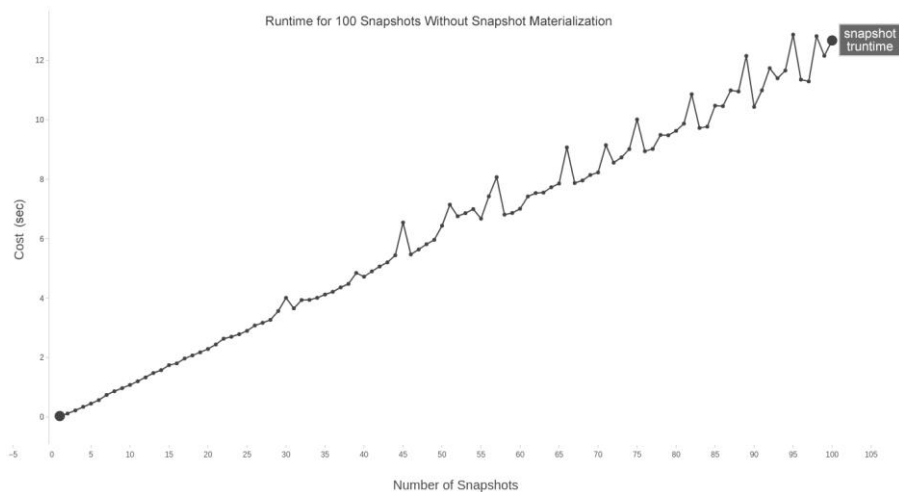
The snapshot creation and blockchain verification take linear time.

$$\mathcal{O}(|\{x : x \in r^T \text{ and } x.\text{updates} \leq t\}|) \simeq \mathcal{O}(t)$$

Recomputing snapshots and blockchain verification is expensive.

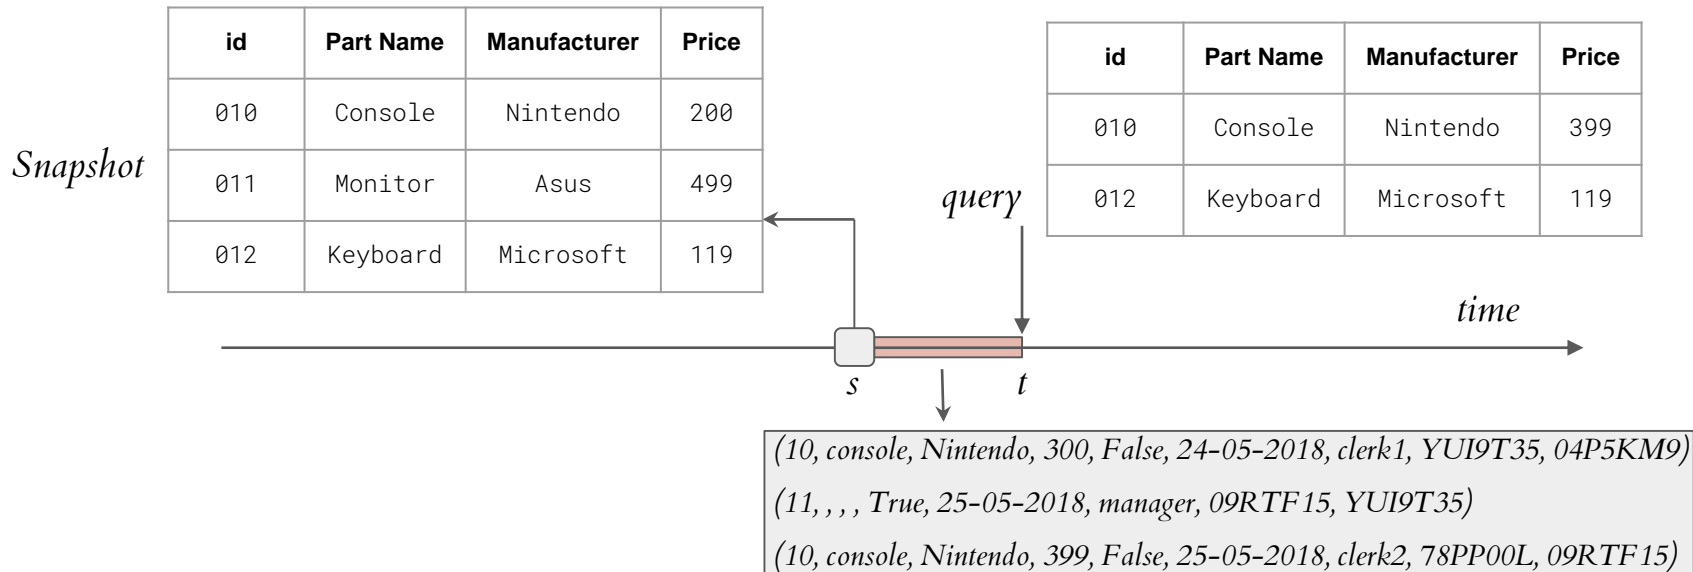
Temporal databases: snapshot creation

Runtime of creating 1-100 snapshots



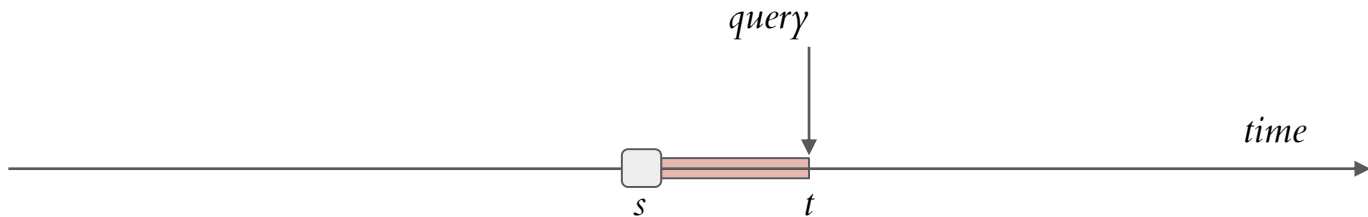
Lowering the cost of snapshot generation

We propose to *precompute* snapshots for *materialization* to lower the cost of snapshot generation.



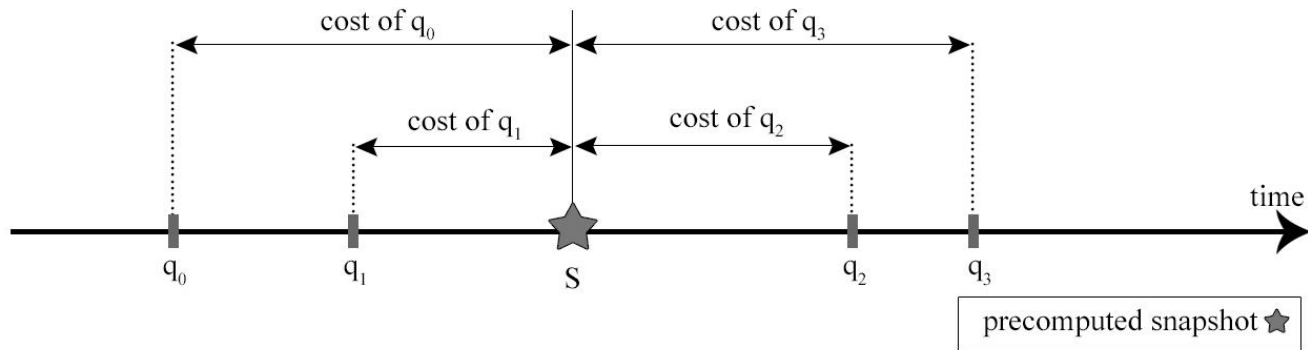
Lowering the cost of snapshot generation

We propose to *precompute* snapshots for *materialization* to lower the cost of snapshot generation.



$$\mathcal{O}(|\{x : x \in r^T \text{ and } x.\text{updates} \in [s, t]\}|) \simeq \mathcal{O}(|s - t|)$$

Overall cost of query using a materialized snapshot



Cost of query using single materialized snapshot is

$$\text{cost}(T_q|s) = \sum_{q \in T_q} |q - s|$$

Lowering the cost of snapshot generation

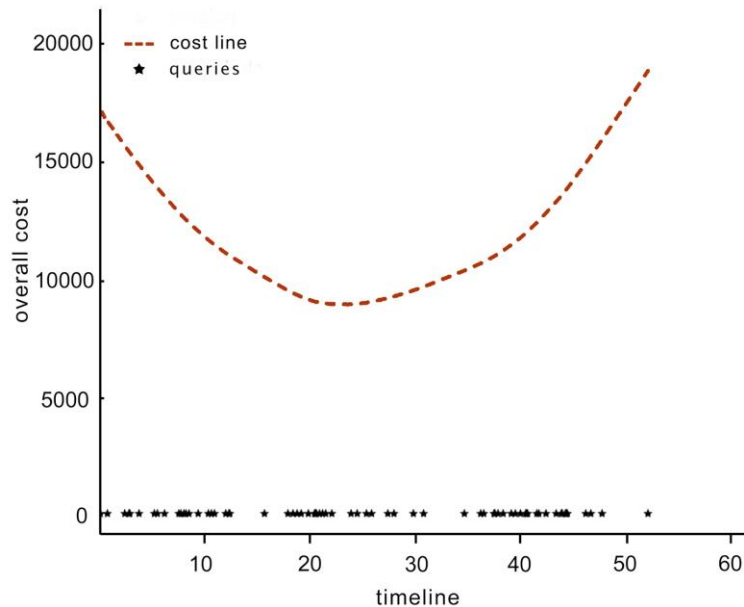
Our goal is to precompute *multiple optimal* snapshots for materialization



$$cost(T_q|S) = \mathit{Argmin}(\sum_{q \in T_q} \{|q - s| : s \in S\})$$

Single snapshot materialization

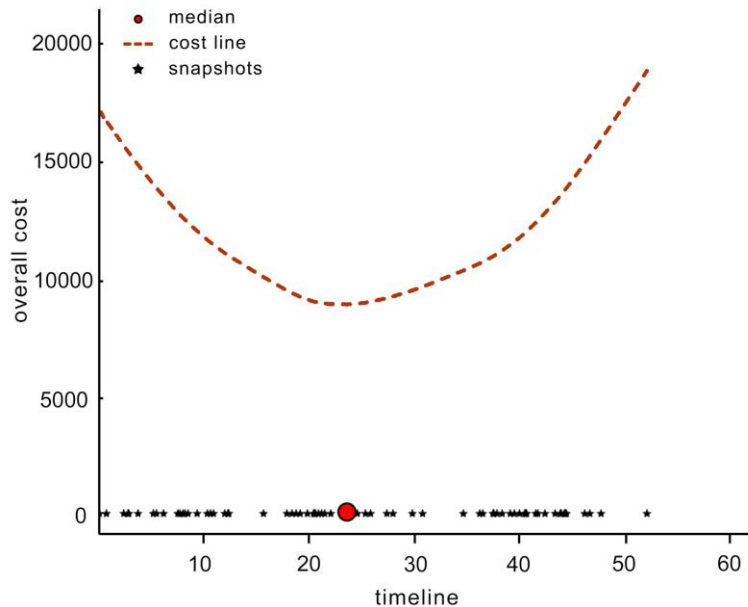
*Where is the optimal timestamp
for snapshot placement?*



Single snapshot materialization

Theorem. The median of queries minimizes the sum of absolute deviation

*If the snapshot was in the median of the performed queries, it was in the **optimal** timestamp.*



Lowering the cost of snapshot generation

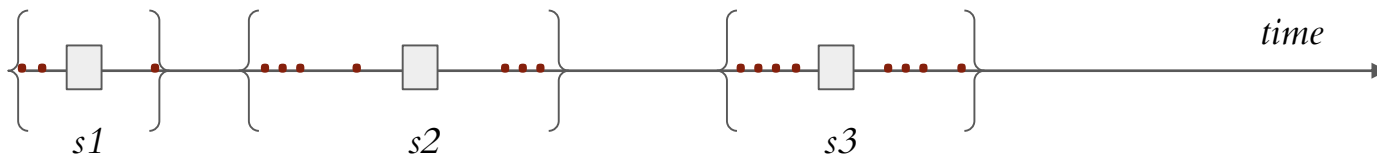
*We suggest to store the
timestamp of previous queries*

*It gives us insight into the pattern of
queries and hotspots of the timeline*



Lowering the cost of snapshot generation

How can we create *optimal segmentations* in which if a snapshot is placed for materialization, the overall cost of query answering is minimum?



Multiple optimal snapshot placement: possible solutions

The segmentations could be computed with:

1. Recursive algorithm
2. Dynamic programming
3. Heuristic method

Which one is better?

Optimal query segmentation: Recursive Algorithm

We try all the last segment of the queries and we choose the one with the lowest cost.

Benefits:

The optimal solution is guaranteed

Drawbacks:

Inefficiency in computation for large number of queries

Function computeOPT(Q, m):

```
 $n = |Q|$   
 $\text{OPT}[i, 0] = \infty$   
for  $k \leftarrow 1$  to  $m$  do  
  for  $i \leftarrow 1$  to  $n$  do  
     $j^* = \underset{j \in [1, i]}{\operatorname{argmin}} (\text{cost}(\text{OPT}[j, k-1]) + \text{cost}(Q[j+1, n]))$   
     $\text{OPT}[i, k] = \text{OPT}[j^*, k-1] \cup \{\text{median}(Q[j+1, n])\}$   
  end  
end
```

The recursive algorithm requires $\mathcal{O}(2^m)$

Optimal query segmentation: Dynamic Programming

The result of expensive function calls are stored in a table, and retrieved when the same input is given to the function.

Benefits:

The optimal solution is guaranteed and has lower cost than recursion.

Drawbacks:

Inefficiency in computation for larger number of queries

Function computeOPT(Q, m):

```
 $n = |Q|$   
 $minVal = \infty$   
for  $i \leftarrow 1$  to  $m$  do  
  for  $j \leftarrow 1$  to  $n + 1$  do  
    for  $k \leftarrow 1$  to  $j$  do  
       $minVal = \min(minVal, Table[i, k] + cost(Q[j - k :]))$   
    end  
     $Table[i, j] = minVal$   
  end  
end
```

The dynamic programming requires $\mathcal{O}(mn^2)$

Optimal query segmentation: Heuristic method

Random timestamps are chosen and partitions are created. Then the random timestamps are moved on the timeline until the within cluster sum of squares is minimized.

Benefits:

Low time complexity

Drawbacks:

The optimal solution not guaranteed.

The K-means clustering algorithm is given as Algorithm 9.

Function K-Means ($T_q^* \{q_1, \dots, q_n\}, m, \text{maxIteration}$):

```
    iteration  $\leftarrow$  0
    repeat
         $\{\mu_1, \dots, \mu_m\} \leftarrow \text{SelectRandomSeeds}(\{q_i \in T_q^*\}, m)$ 
        for  $i \leftarrow 1$  to  $n$  do
             $J \leftarrow \text{argmin}_{j^*} \|\mu_{j^*} - q_i\|^2$ 
             $\mathcal{L}_j \leftarrow \mathcal{L}_j \cup \{q_i\}$ 
        end
        for  $j \leftarrow 1$  to  $m$  do
             $\mu_j \leftarrow \frac{1}{|\mathcal{L}_j|} \sum_{q \in \mathcal{L}_j} q$ 
        end
        iteration ++
    until convergence and iteration  $\leq$  maxIteration
    return  $\{\mu_1, \dots, \mu_m\}$ 
```

The K-Means clustering requires $\mathcal{O}(m * k * n)$

Performance of optimal snapshot computation: experiment

Experiment 1:

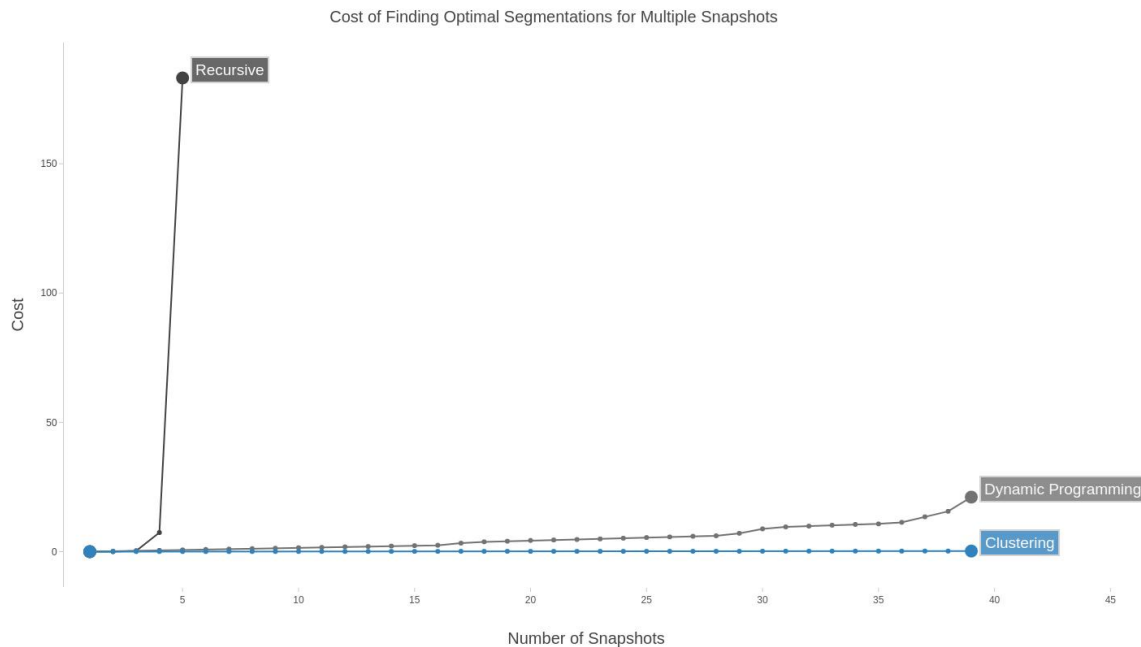
Fixed number of queries: 140

variable number of snapshots: 1 - 40

Objective:

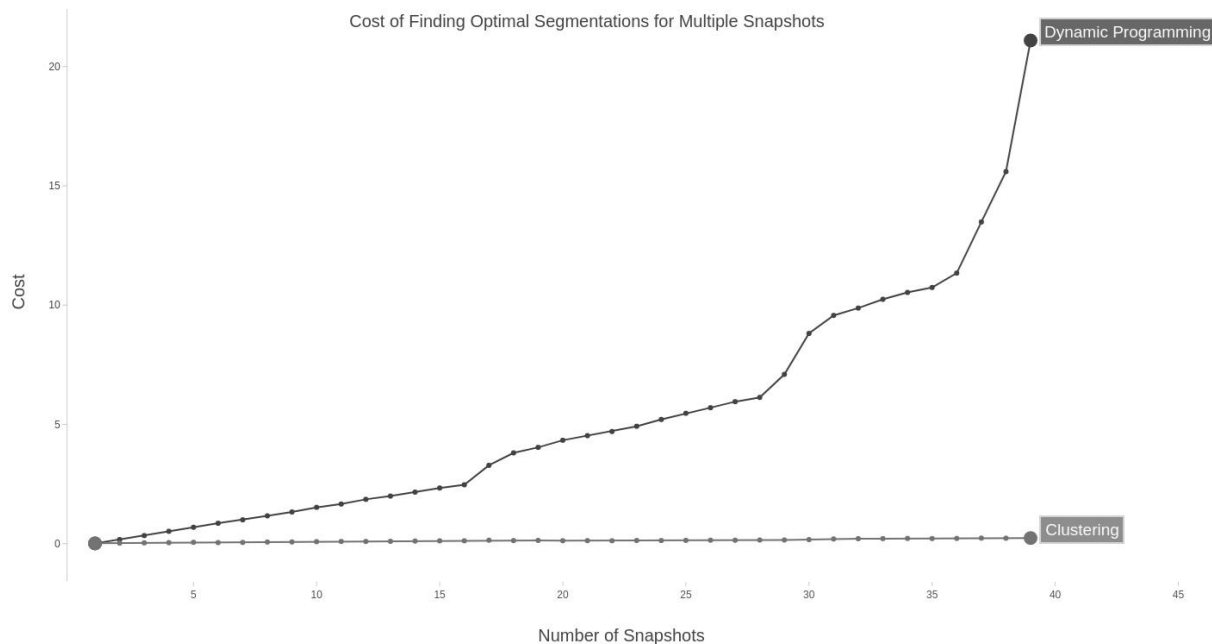
Evaluate performance of algorithms with respect to our project

performance of optimal snapshot computation :comparison



Snapshots	Recursive	Dynamic	Clustering
1	0.0002	0.01	0.01
2	0.01	0.17	0.02
3	0.32	0.34	0.03
4	7.39	0.52	0.04
5	183.18	0.69	0.06
10	N/A	1.52	0.08
15	N/A	2.33	0.12
20	N/A	4.33	0.12
25	N/A	5.46	0.14
30	N/A	8.81	0.17
35	N/A	10.74	0.21
40	N/A	21.09	0.24

Performance of optimal snapshot computation :comparison

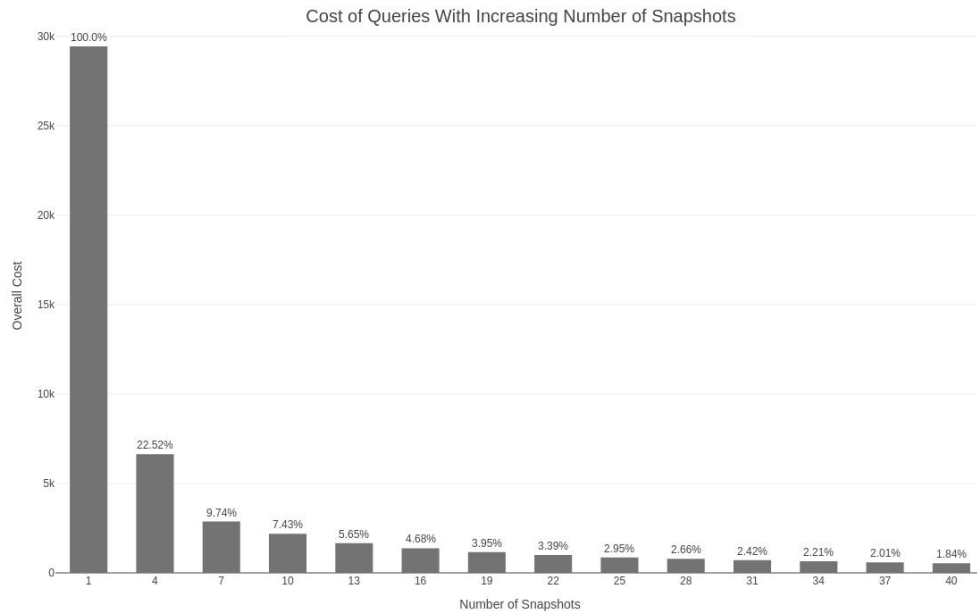


Snapshots	Dynamic	Clustering
1	0.01	0.01
2	0.17	0.02
3	0.34	0.03
4	0.52	0.04
5	0.69	0.06
10	1.52	0.08
15	2.33	0.12
20	4.33	0.12
25	5.46	0.14
30	8.81	0.17
35	10.74	0.21
40	21.09	0.24

Experiment 1. Effectiveness of snapshot materialization

Case: The optimal segmentations computed using dynamic programming

Objective: To see how increasing number of materialized snapshot affects the cost



Performance of optimal snapshot computation: experiment

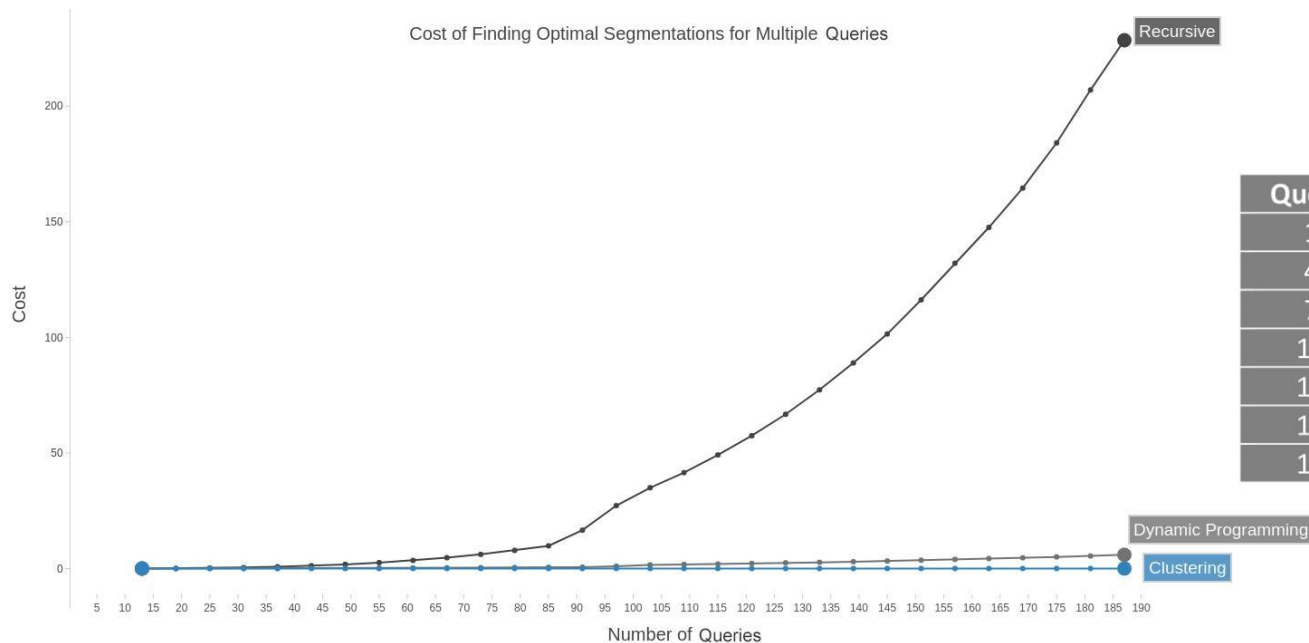
Experiment 2:

Fixed number of snapshots: 4
variable number of queries: 12 - 190

Objective:

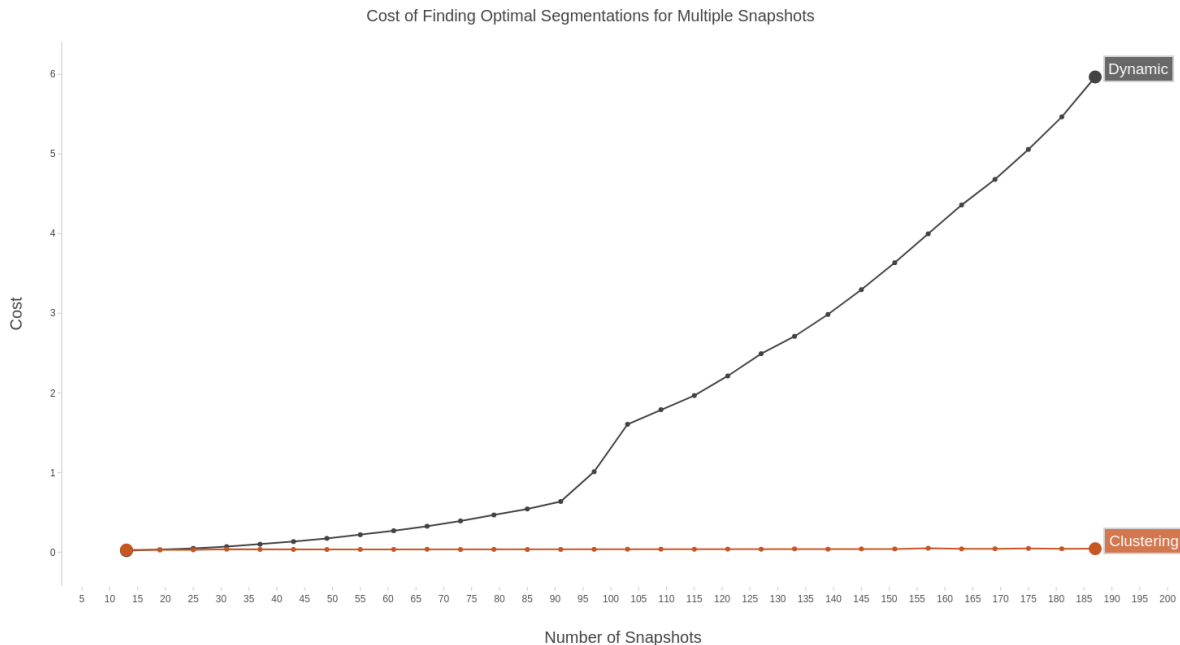
Evaluate performance of algorithms with
respect to our project

Performance of optimal snapshot computation: comparison



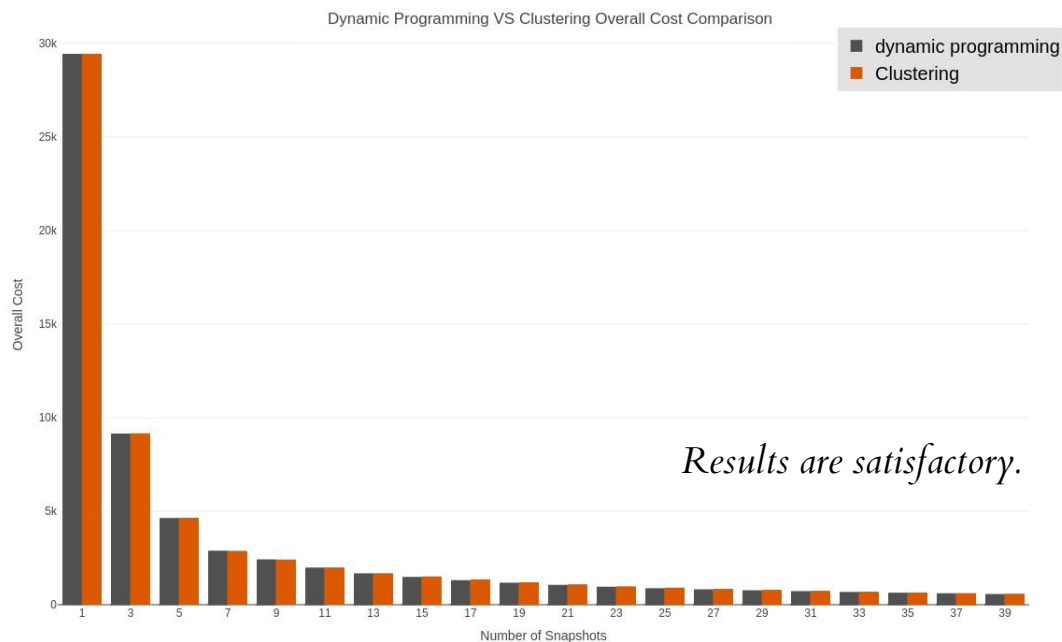
Queries	Recursive	Dynamic	Clustering
13	0.05	0.01	0.02
43	1.23	0.13	0.03
73	6.15	0.39	0.03
103	34.98	1.60	0.03
133	77.31	2.71	0.04
163	147.57	4.35	0.04
187	228.52	5.96	0.04

Performance of optimal snapshot computation: comparison



Queries	Dynamic	Clustering
13	0.01	0.02
43	0.13	0.03
73	0.39	0.03
103	1.60	0.03
133	2.71	0.04
163	4.35	0.04
187	5.96	0.04

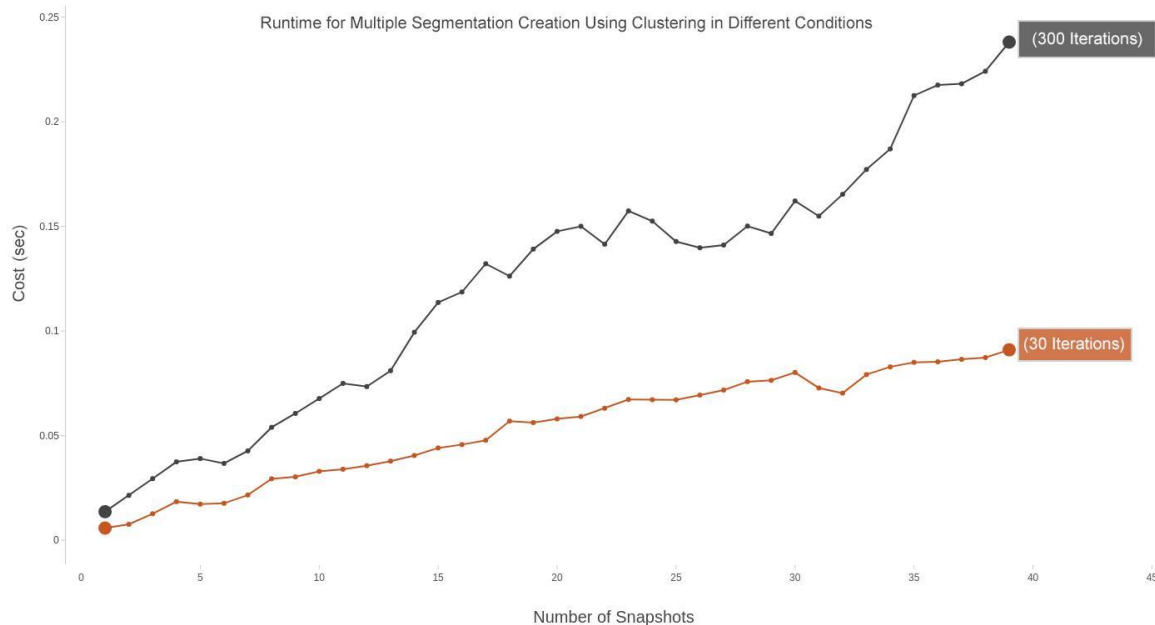
Performance of optimal snapshot computation: comparison



Snapshots	Dynamic	Clustering
1	29439.26	29439.26
3	9141.55	9159.13
5	4626.77	4630.08
7	2867.74	2867.74
9	2410.46	2412.62
11	1972.14	1980.95
13	1664.14	1673.32
15	1471.58	1509.57
17	1300.32	1351.44
19	1162.25	1194.61
21	1051.97	1079.711
23	951.06	970.72
25	867.35	907.30
27	810.01	836.97
29	759.69	787.86
31	713.04	731.61
33	670.39	684.13
35	629.69	640.58
37	591.08	608.96
39	557.81	575.97

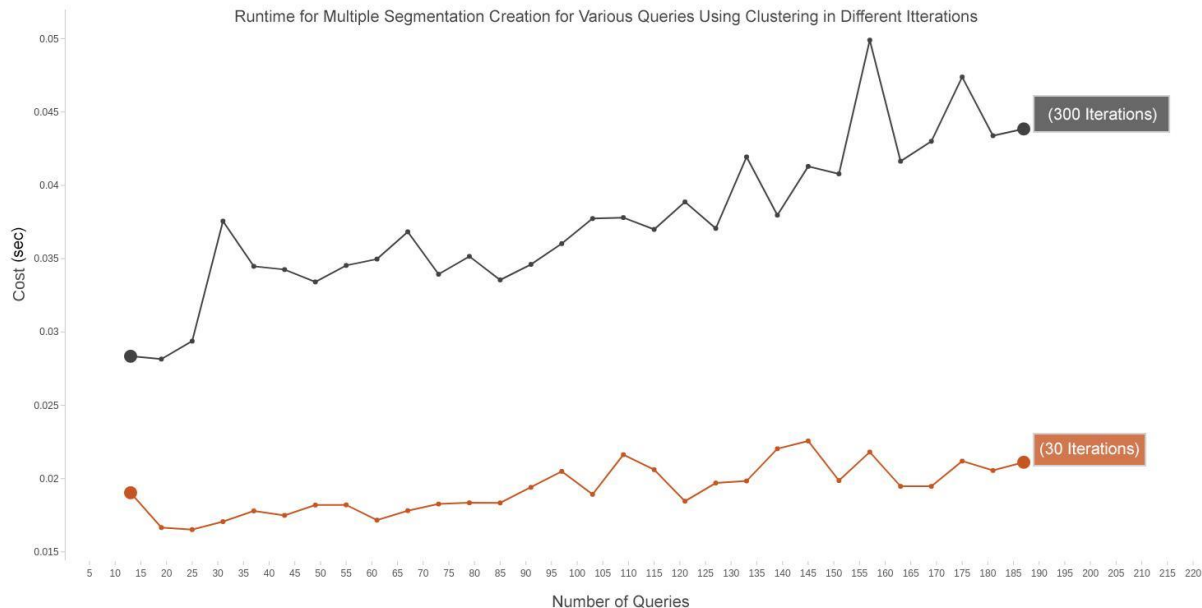
Performance of optimal snapshot computation: comparison

Runtime for variable number of snapshots but fixed number of queries



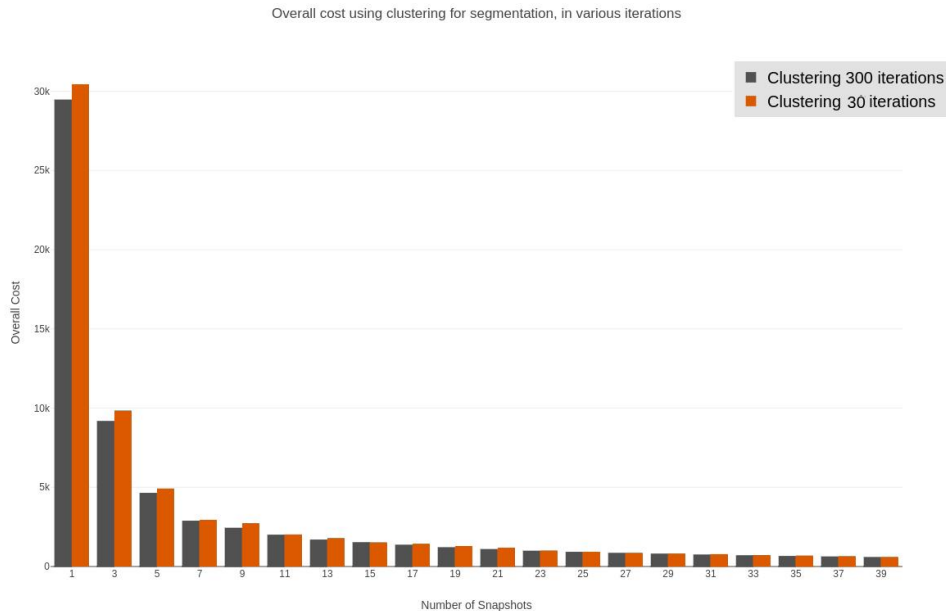
Performance of optimal snapshot computation: comparison

Runtime for fixed number of queries but variable number of snapshots



Performance of optimal snapshot computation: comparison

Comparison between the overall cost of query answering using optimal segmentations made by heuristic method in different iterations



Performance of optimal snapshot computation: experiment

Experiment 3:

40 snapshots with multiple segmentation strategy
on a timeline with 1000 queries.

Objective:

To see if finding optimal segmentations worth it.

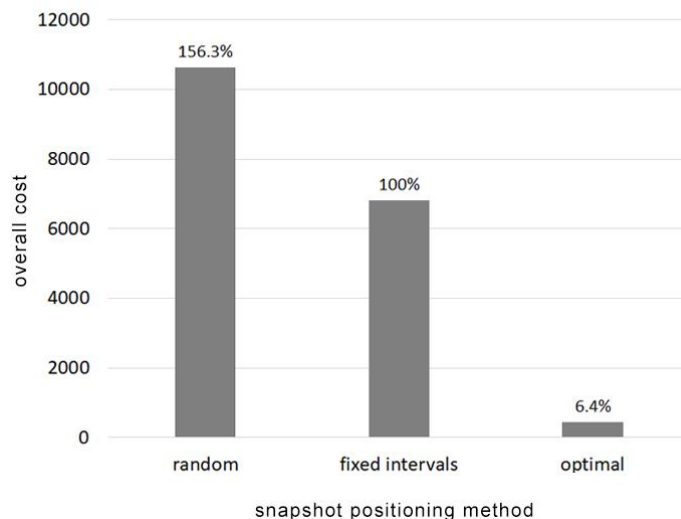
Experiment 1. Effectiveness of snapshot materialization

Case 1: Segmentations created randomly

Case 2: Fixed segmentations were created

Case 3: Optimal segmentation was computed using dynamic programming

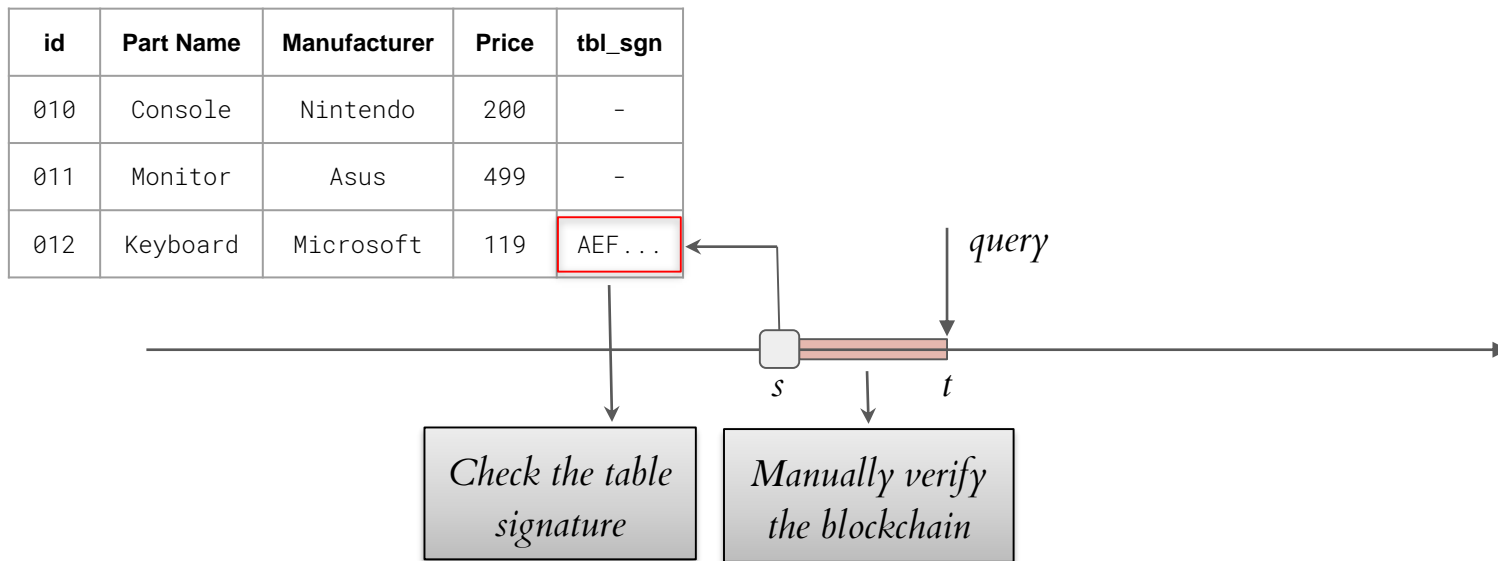
Objective: Does optimal segmentation computation make any difference? YES



Further remarks

Lowering the cost of blockchain verification

We propose to *digitally sign* the *precompute* snapshots.

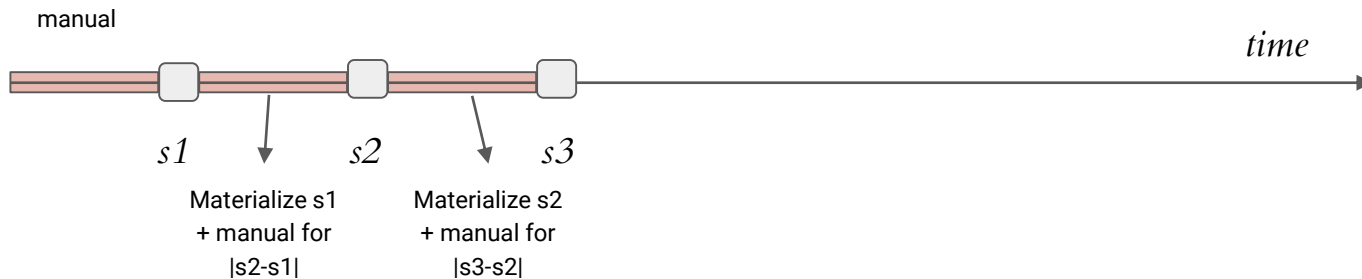


Materialized snapshot generation

How can we reduce the cost of snapshot generation?

For the first snapshot, computation is done manually.

For the subsequent snapshots, previous snapshots are materialized.

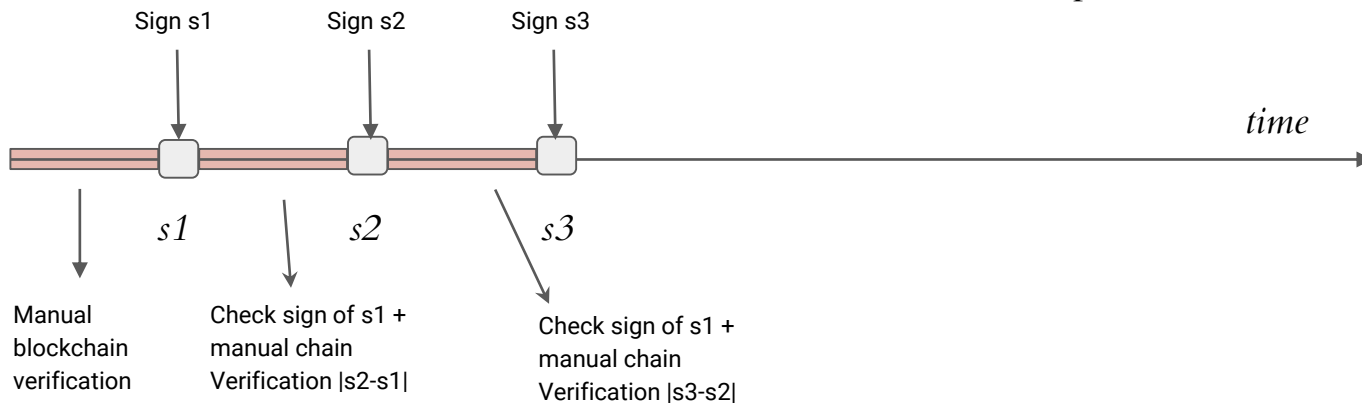


Materialized Blockchain verification

How can we reduce the cost of snapshot signing?

For the first snapshot, Blockchain verification is done manually.

For the subsequent snapshots, previous snapshots are materialized.



Conclusion

- Using the temporal relation that utilizes Blockchain we can provide trusted data provenance.
- Snapshot materialization has promising results in reducing the cost of query answering.
- With storing the timestamp of performed queries we can find the hotspots on the timeline and we can create optimal segmentations that materialize their exclusive snapshot
- Optimal snapshot placement has more impact in reducing cost than random or fixed interval snapshot placement.

Conclusion

- Dynamic programming, recursion and heuristic methods could be used for optimal segmentation computation.
- In this project, the heuristic method is favorable because of efficient cost and satisfactory results.
- To reduce the cost of multiple snapshot placement as well as chain verification, previous snapshots could be materialized.
- Snapshots could be dynamically repositioned as the new queries performed on the system.

Towards Trusted Temporal Databases Using Blockchain

By Amin Beirami
MSc student, UOIT

Supervisors: Dr. Ken Pu and Dr. Ying Zhu

October 2018